

**Exercise 9 (Serializability)** Consider the following transactions:

- T1: read(A,t); t:=t+2; write(A,t); read(B,t); t:=t\*3; write(B,t)
- T2: read(B,s); s:=s\*2; write(B,s); read(A,s); s:=s+3; write(A,s)

As usual, we assume that whatever consistency constraints are given, these are preserved when T1 and T2 execute in isolation.

- a) Show that for the above T1, T2 the schedules (T1;T2) and (T2;T1) have the same effect on the database (this is usually *not* the case).

$result(T1;T2) : A_1 := A_0 + 2; B_1 := 3 * B_0; B_2 := 2 * B_1; A_2 := A_1 + 3$   
 $\Rightarrow A_2 = A_0 + 5; B_2 = 6 * B_0; \text{ the same result state } (A_2, B_2) \text{ is reached for } (T2;T1).$

- b) Give a serializable and a non-serializable schedule for T1 and T2.

*T1 modifies A, then B; T2 the other way round. The serial schedules above are trivially serializable. A really interleaved serializable schedule is  $S = read_1(A, t); read_2(B, s); t_1 := t_1 + 2; s_2 := s_2 * 2; write_2(B, s); write_1(A, t); read_2(A, s); read_1(B, t); t_1 := t_1 * 3; write_1(B, t); s_2 := s_2 + 3; write_2(A, s)$ . But note that  $S$  is not conflict serializable! A non-serializable schedule is obtained, e.g., by moving  $T2$ 's update on  $A$  between the read and write on  $A$  of  $T1$ . Then  $T2$ 's effect on  $A$  is lost!*

- c) How many serial schedules are there?

*A schedule is **serial** if for any two actions  $p_i$  from  $T_i$  and  $p_j$  from  $T_j$  follows: if  $p_i <_S p_j$  then all  $T_1$  actions are before all  $T_j$  actions. Otherwise it is **interleaved**. For two transactions, there are two serial schedules (T1;T2) and (T2;T1).*

- d) How many serializable schedules are there?

*NB: Here we do **not** ask for conflict serializable schedules, but just serializable schedules (i.e., which will always produce the same result as some serial schedule). In this example,  $S$  is **not** serializable, if a transaction  $T'$  reads the old value ( $A$  or  $B$ ) after  $T$  has read but before  $T$  has written the new value. Then one update will be lost. E.g., we cannot move the  $r(A)$  of  $T2$  before the  $w(A)$  but after the read of  $T1$ . However the first 3 actions of  $T1$  and  $T2$  can be freely interleaved, similarly, the last 3 actions. So we get  $\binom{6}{3}\binom{6}{3} + 2 = 402$  serializable schedules (including the two serial schedules); cf. Exercise 11.*

**Exercise 10 (Serializability)** Consider the schedule  $S =$

$r1(A); r2(A); r3(B); w1(A); r2(C); r2(B); w2(B); w1(C)$

- a) Give the precedence graph  $P(S)$ .

$T3 \rightarrow T2 \rightarrow T1$

- b) Is the schedule conflict-serializable? If so, give all equivalent serial schedules.

*Yes, since  $P(S)$  is acyclic. The only equivalent serial schedule is (T3; T2; T1).*

**Exercise 11 (Interleavings)** Given two transaction  $T1$  and  $T2$  consisting of  $n1$  and  $n2$  actions. How many interleavings (i.e., different schedules with  $T1$  and  $T2$ ) are possible?

*Any schedule for  $T1$  and  $T2$  has  $n1+n2$  "positions";  $n1$  of those will be used for  $T1$  actions: we can choose those  $n1$  arbitrarily from all positions; after that everything is fixed. Hence there will be  $\binom{n1+n2}{n1} = \binom{n1+n2}{n2} = \frac{(n1+n2)!}{n1!n2!}$  of those.*

**Exercise 12 (Serializability, s.154)** Show that the converse of the Lemma does not hold, i.e.,  $P(S1) = P(S2)$  does *not* imply that  $S1$  and  $S2$  are conflict equivalent.

*Consider  $S1 = w1(A); r2(A); w2(B); r1(B)$  and  $S2 = r2(A); w1(A); r1(B); w2(B)$ . Then  $P(S1) = P(S2) = \{T1 \rightarrow T2 \rightarrow T1\}$  but we cannot swap non-conflicting actions to obtain one schedule from the other.*

**Exercise 13 (Deadlocks, Starvation, s.163ff)** We speak of *starvation* if a transaction is repeatedly rolled back and never finishes. Can the following anti-deadlock strategies lead to starvation (explain/give an example):

- transaction which have reached their time limit are rolled back

*Can lead to starvation: e.g., if the limit is too short even for a non-deadlocked execution.*

- transactions which would introduce a cycle in the waits-for graph are rolled back

*Can lead to starvation: after roll back, a transaction may again become the one that introduces a cyclic wait.*

- lock request can appear only in a predefined order

*Cyclic wait conditions are prevented in the first place, so transactions are never rolled back. Can they wait infinitely? Not if we have a “fair” scheduler, which allows transactions to make progress eventually. For example, if we distribute available locks on a first-come first-served basis, starvation is prevented.*

- transactions are rolled back based on their timestamps (WAIT-DIE or WOUND-WAIT scheme)

*Starvation is prevented since eventually any transaction will become oldest and thus cannot be rolled back.*

**Exercise 14 (Deadlocks, s.167)** Show that the WAIT-DIE and WOUND-WAIT strategies prevent deadlocks.

*A deadlock is a cycle in the waits-for graph. However, for each of the schemes, edges only go in one direction (either wait for older or wait for younger). Hence there cannot be cycles.*

**Exercise 15 (Serializability)** Consider the following two transactions:

- $T_0$ :  $r(A); r(B); \text{if } A=0 \text{ then } B:=B+1; w(B);$
- $T_1$ :  $r(B); r(A); \text{if } B=0 \text{ then } A:=A+1; w(A);$

Assume initial values  $A=B=0$ , and that we have the consistency requirement:  $A=0 \text{ OR } B=0$ .

- Show that every serial execution maintains consistency.
- Give a nonserializable schedule.
- Is there a serializable, non-serial schedule?

**Exercise 16 (2PL, s.162)** Show that the schedule in s.162 is

- serializable and
- explain why it cannot be obtained from a 2PL scheduler.

**Exercise 17 (View/Conflict Serializability)** Consider the following schedule  $S$ :

$T_1$	$T_2$	$T_3$
$w(A)$	$r(A)$	
		$w(B)$
$w(B)$		$w(B)$
	$w(A)$	$r(B)$
	$r(B)$	

- Give the precedence graph  $P(S)$ .
- Is  $S$  conflict serializable? Explain.
- Is  $S$  view serializable? Explain.

$P(S) = T_1 \rightarrow T_2, T_1 \rightarrow T_3, T_3 \rightarrow T_1$ , i.e., cyclic, so not conflict serializable.					
	$T_b$	$T_1$	$T_2$	$T_3$	$T_f$
	$w(A)$				
	$w(B)$				
		$w(A)$	$r(A)$		
		$w(B)$		$w(B)$	
			$w(A)$	$w(B)$	
			$r(B)$	$r(B)$	
					$r(A)$
					$r(B)$
For view serializability, add $T_b$ and $T_f$ :					and construct
the labeled precedence graph. The crucial edges are $T_1 \xrightarrow{B,1} T_3$ and $T_2 \xrightarrow{B,1} T_1$ . By picking the former instead of the latter, we get an acyclic graph; its topological sort yields the serial schedule $(T_1; T_3; T_2)$					

**Exercise 18 (Serializability)** Two transactions are *not interleaved* in a schedule  $S$  if every operation of one transaction precedes every operation of the other in  $S$ . (Note:  $S$  need not be serial.) Give an example of a *serializable* schedule  $S$  with the following properties:

- $T_1$  and  $T_2$  are not interleaved in  $S$ , and
- $T_1$  precedes  $T_2$  in  $S$ , and
- in any serial schedule equivalent to  $S$ ,  $T_2$  precedes  $T_1$ .

Hint:  $S$  may have additional transactions.

	$T_1$	$T_2$	$T_3$
			$w(A)$
Consider this schedule:	$r(A)$		
		$r(B)$	
			$w(B)$
In an equivalent serial schedule we have to make sure that $T_3 < T_1$ (because of $w_2(A), r_1(A)$ ) and that $T_2 < T_3$ (because of $r_2(B), w_3(B)$ ). Hence we have $(T_2; T_3; T_1)$ .			

**Exercise 19 (IO Cost, s.115ff)** Consider two relations R and S with  $T(R) = 1000$ ,  $T(S) = 500$ , and  $S(R) = S(S)$  such that 10 tuples fit into a block. Determine the number of disk IOs for the following joins:

- a) nested-loop join with (i) non-clustered, tuple access, (ii) non-clustered, block access, (iii) clustered, block access,
- b) merge join with clustered relations (i) given R, S sorted, (ii) R, S not sorted,
- c) index join: consider the different expected number of matching R tuples from s.129, i.e., 1, 2, and 0.01
- d) hash join with ideal distribution across buckets.

**Exercise 20 (Join Algorithms)** Which are typical “good uses” of nested-loop join, hash join, merge join, and index join algorithms?

*nested loop: small relations only; hash: equi-joins on unsorted, non-indexed relations; merge: sorted relations and non-equi joins like  $R.A > S.B$ ; index: if index available, but consider expected number of matches*

**Exercise 21 (Recovery)**

- a) Undo logging requires that before an item is modified on disk, the log records pertaining to X are on disk (WAL: write ahead logging). Show using an example that an inconsistent database may result if log records for X are not flushed to disk before X.
- b) Show using an example that an inconsistent database may result if some items are not written to disk before the commit is written to the log (even if WAL holds).
- c) Redo logging: show using an example that an inconsistent database may result if some items are written to disk before the commit is written on the log (even if WAL holds)