

# Concurrency Control / Serializability Theory

**Correctness of a program can be characterized by invariants**

- **Invariants over state**

Linked list:

For all items I where I.prev != NULL: I.prev.next == I

For all items I where I.next != NULL: I.next.prev == I

- **Invariants over operations**

Stack:

Stack.size after push == Stack.size before push + 1

Stack.size after pop == Max (Stack.size before pop - 1, 0)

**Consistency can be defined as validity of invariants**

*Non atomic actions temporarily violate invariants, hence program moves from consistent to consistent state through potentially inconsistent states*

**Running programs in parallel means inconsistent states can be observed and acted upon, which is what we need to prevent. This is the purpose of concurrency control.**

# Modeling Transactions

**Transaction is a sequence of read and write actions on data ...**

- **A transaction  $T_i$  is a partial order with ordering relation  $<_i$  where**
  - $T_i \subseteq \{ r_i[x], w_i[x] \mid x \text{ is a data item} \} \cup \{ a_i, c_i \}$  ;
  - $a_i \in T_i$  iff  $c_i \notin T_i$ ;
  - if  $t$  is  $c_i$  or  $a_i$ ,  $t \in T_i$ , for any other operation  $p \in T_i$ ,  $p < t$  ;
  - if  $r_i[x], w_i[x] \in T_i$  then either  $r_i[x] < w_i[x]$  or  $w_i[x] < r_i[x]$  .
- **Uninterpreted features – we do not make any assumptions**
- **Transactions are drawn as direct acyclic graphs (DAGs)**

# Execution History

Let's model execution history from transactions, but we need to define conflicting operations

- **Conflicting operations operate upon the same data and at least one of them is write**
- **Conflicting transactions contain conflicting operations**

**Execution history**

- **$T = \{T_1, T_2, \dots, T_n\}$  is a set of transactions**
- **Complete history is a partial order with ordering relation  $<_H$  where**
  1.  **$H = \bigcup_n^{i=1} T_i$**
  2.  **$<_H \supseteq \bigcup_n^{i=1} <_i$**
  3. **for any two conflicting operations  $p, q \in H$ , either  $p <_H q$  or  $q <_H p$**
- **History is simply a prefix of complete history**

*As well as with particular transactions, we draw history as a DAG*

*We don't draw all arrows implied by transitivity*

**Committed, aborted, and active transactions in history H**

- **Committed projection  $C(H)$  of H contains only operations of committed transactions**

# Serializable Histories

**History is serializable if it is equivalent to a serial history**

- **H and H' equivalent if**
  - **they have the same transactions and operations**
  - **they order conflicting operations of non-aborted transactions in the same way**




*The outcome of concurrent execution of transactions depends  
on the relative ordering of conflicting operations  
Executing two non-conflicting operations in either order  
has the same computational effect*

- **Complete history H is serial if for every two transactions  $T_i, T_j \in H$  all operations of  $T_i$  are before operations of  $T_j$  or vice versa**

*We require completeness of history since incomplete  
transactions do not preserve data consistency*

- **History is serializable (SR) if its committed projection is equivalent to a serial history**

*Only execution of committed transactions is guaranteed*



# Serializability Theorem

## How we recognize serializable history

- For  $H$  a history over  $T = \{T_1, T_2, \dots, T_n\}$ , serialization graph  $SG(H)$  is directed graph
  - nodes are transactions in  $T$  that are committed in  $H$
  - edges are  $T_i \rightarrow T_j$  if one of  $T_i$ 's operations conflict and precedes with one of  $T_j$ 's operations in  $H$

*Single edge for more than one pair of conflicting operations*

*No transitivity*

*Edges between transactions implies ordering in a serial history*

*We can find an equivalent serial history if  $SG(H)$  is acyclic*

## Serializability theorem

- A history  $H$  is serializable iff  $SG(H)$  is acyclic

$$\begin{array}{rcccl}
 & & & r_3[x] \rightarrow w_3[x] \rightarrow c_3 & \\
 & & \nearrow & & \\
 H_\varsigma = & r_1[x] \rightarrow w_1[x] \rightarrow w_1[y] \rightarrow c_1 & & & \\
 & \nearrow & \nearrow & & \\
 & r_2[x] \rightarrow w_2[y] \rightarrow c_2 & & & 
 \end{array}$$

$$\text{SG}(H_\varsigma) = \quad T_2 \xrightarrow{\quad} T_1 \rightarrow T_3$$

# Recoverable Histories

To ensure correctness in the presence of failures, the scheduler must produce SR histories that are recoverable

- $T_i$  reads  $x$  from  $T_j$  if
  - $w_j[x] < r_i[x]$
  - $a_j$  does not precede  $r_i[x]$  in the partial order
  - if there is some  $w_k[x]$  such that  $w_j[x] < w_k[x] < r_i[x]$  then  $a_k < r_i[x]$

- $T_i$  reads from  $T_j$  if there is an  $x$  such that  $T_i$  reads  $x$  from  $T_j$

*A transaction can read a data item from itself*

- $H$  is recoverable (RC) if, whenever  $T_i$  reads from  $T_j$  ( $i \neq j$ ) in  $H$  and  $c_i \in H$ , then  $c_j < c_i$

*Intuitively, if  $T_i$  reads  $x$  from  $T_j$ , it has to wait whether  $x$  will not be invalidated by  $a_j$*

- $H$  avoids cascading aborts (ACA) if, whenever  $T_i$  reads  $x$  from  $T_j$  ( $i \neq j$ ) in  $H$ , then  $c_j < r_i[x]$

*A transaction may read only those values that are written by committed transactions or by itself*



- **H is strict (ST) if, whenever  $w_j[x] < o_i[x]$  ( $i \neq j$ ), either  $a_j < o_i[x]$  or  $c_j < o_i[x]$  where  $o_i[x]$  is  $r_i[x]$  or  $w_i[x]$**

*No data item may be read or overwritten until the transaction that previously wrote into it terminates*

- **Examples**

**T1 = w1[x] w1[y] w1[z] c1      T2 = r2[u] w2[x] r2[y] w2[y] c2**

**H1 = w1[x] w1[y] r2[u] w2[x] r2[y] w2[y] c2 w1[z] c1**

**H2 = w1[x] w1[y] r2[u] w2[x] r2[y] w2[y] w1[z] c1 c2**

**H3 = w1[x] w1[y] r2[u] w2[x] w1[z] c1 r2[y] w2[y] c2**

**H4 = w1[x] w1[y] r2[u] w1[x] c1 w2[x] r2[y] w2[y] c2**


**H1 is not RC (T2 reads y from T1 and  $c2 < c1$  )**

**H2 is RC but not ACA (T2 reads y from T1 before T1 is committed)**

**H3 is ACA but not ST (T2 overwrites value written to x by T1 before T1 terminates)**

**H4 is ST**

- **$ST \subset ACA \subset RC$**



## Prefix Commit-Closed Properties

**If a scheduler produces a “correct” history  $H$ , then any prefix  $H'$  of  $H$  should be also correct with respect to committed transactions**

*In case of DBS fail after  $H'$ , the database reflects  $C(H')$  which should be correct*

- **A property is prefix commit-closed if, whenever the property is true of history  $H$ , it is also true of history  $C(H')$ , for any prefix  $H'$  of  $H$** 
  - **SR is prefix commit-closed**
  - **RC, ACA, and ST are prefix commit-closed**

# Operations Beyond Reads and Writes

**Let's consider other set of operations**

- **New definition of conflict**
  - **Compatibility matrix for all operations**
- **SG definition unchanged**
- **Serialization theorem unchanged**

*We use read and write operations for simplicity and practical usability in databases  
Serializability theory works correctly with different operations*

# View Equivalence

**Let's consider other criterion for history equivalence**

- **We don't know what computation  $f(x)$  is between  $r[x]$  and  $w[x]$** 
  - **We know that it is some function of all reads**
  - **Thus, if all reads are the same in two histories, then all writes should also be the same**

**Consequence:**

- **If each transaction reads each of its data items from the same writes in both histories, then all writes write the same values in both histories**
- **If for each  $x$ , the final  $w[x]$  is the same in both histories, then the final value of all data is the same in both histories**

*If both histories leave database in the same final state,  
then the histories have to be considered equivalent*

# View Equivalence

**More formally:**

- **Final write of  $x$  in  $H$  is  $w_i[x] \in H$ , such that  $a_i \notin H$  and for any  $w_j[x] \in H$  ( $j \neq i$ ) either  $w_j[x] < w_i[x]$  or  $a_j \in H$**
- **Two histories  $H, H'$  are equivalent if**
  - **they are over the same set of transactions and have the same operations**
  - **for any  $T_i, T_j$  such that  $a_i, a_j \notin H$  (hence  $a_i, a_j \notin H'$ ) and for any  $x$ , if  $T_i$  reads  $x$  from  $T_j$  in  $H$  then  $T_i$  reads  $x$  from  $T_j$  in  $H'$**
  - **for each  $x$  if  $w_i[x]$  is the final write of  $x$  in  $H$  then it is also the final write of  $x$  in  $H'$**

**Call this equivalence view equivalence**

- **Used for concurrency control algorithms for multicopy data**
  - **Multiversion concurrency control**
  - **Replicated data concurrency control**

**Call the old definition conflict equivalence (conflicting operations of unaborting transactions appear in the same order on both histories)**

# View Serializability

## Recall Conflict Serializability (CSR)

- **H is (conflict) serializable if its committed projection  $C(H)$  is (conflict) equivalent to some serial history**

## View Serializability (VSR)

- **H is view serializable if for any prefix  $H'$  of H,  $C(H')$  is view equivalent to some serial history**
  - **We should emphasize “for any prefix” to ensure prefix commit-closed property**

## VSR is a (strictly) more inclusive concept than CSR

- **If H is CSR then it is VSR. The converse is not, generally, true**

## All practical concurrency control algorithms are conflict-based

*An efficient scheduler that produces exactly the set of all view serializable histories can exist only if  $P=NP$   
(This would imply that a wide variety of notoriously difficult combinatorial problems would be solvable by efficient algorithms ☺)*